

BEOPS — pre-paper v2: the thesis, the evidence for it, and what the paper should be

Status: current · Date: 2026-09-09 (evening) · **Supersedes** PRE_PAPER_CONFERENCE_2026-09-09.md as the discussion entry for the paper. It does not supersede the working document, the registry or any evidence; the earlier pre-paper stands unedited as the morning's version.

Authors: prof. dr Darinka Golubović Matić · doc. dr Semir Poturak — **authors of a scientific paper** prepared for *Creating sustainable commUNiTy*, 6th International Conference on Sustainable Environment and Technologies, University Union – Nikola Tesla, 25 September 2026, online (≤ 6 pages A4, abstract ≤ 200 words, APA). They teach at that university; the work is theirs, runs on their own equipment, is non-commercial, and does not speak for the institution or bind it. With **Svemir** (the authors' local AI infrastructure; on this document carried by Claude, Anthropic) as a verified contributor, not an author.

Every figure below was read from a file on the authors' machine or from a page actually opened, on 9 September 2026 between 13:52 and 15:05 UTC. Unverified items are marked as such.

1 · What changed since the morning version, and why it matters

The morning pre-paper described an inventory and offered three possible papers. It was missing a **thesis**: a reason this work should exist rather than another sensor deployment. That thesis now exists, it is argued from sources that are neither ours nor friendly, and it changes the recommended paper from a description into an argument.

The thesis. *A city does not become knowable by acquiring more instruments. The binding constraint is not sensing but the defensible record — and that is what is missing, and what can be built and shown for one city, in public.*

2 · The evidence for the thesis (three anchors, all verified at the primary source)

2.1 The largest comparable deployment says the bottleneck is elsewhere. The Newcastle Urban Observatory operates, on its own account, more than 3,600 sensor streams across 65 variables, ingesting on the order of 10,000 observations per minute, with over ten billion observations collected. Its own team writes: "*Few people want data, what they really want is analysis and insight,*" and that this analysis remains "manual, time consuming, and expensive." The paper also reports what rarely reaches a dashboard: a sensor developing a fault and reporting excessive values; air sensors classified as indicative and permitted to sit within 50 % of a reference; and maintenance becoming a substantial burden past a certain number of units (James et al., 2022, *Frontiers in Sustainable Cities*, 4, 767942, CC BY).

2.2 What the presentation layer routinely omits is exactly the metadata layer. Kitchin and McArdle's six key issues include *veracity and validity* — "data quality metadata is rarely provided" while dashboards "demand trust without transparency" — and *scope and access*, where their own Dublin audit found only **10 of 37** sustainability indicators available at an appropriate scale (Kitchin & McArdle, *Urban data and city*

dashboards: six key issues, Programmable City WP 21; Kitchin, Lauriault & McArdle, 2015, *Regional Studies, Regional Science*, 2(1), doi:10.1080/21681376.2014.983149).

2.3 The new risk in the AI step is now named in the same literature. Huang (2025) defines *geo-hallucination* as model output that is "spatially incoherent or fabricated, despite appearing authoritative", arising from the model's reasoning rather than from measurement, and calls it an epistemic and socio-political risk for planning, governance and equity; the remedy proposed is grounding, critical literacy and a renewed commitment to evidence-based approaches with human oversight (*Environment and Planning B: Urban Analytics and City Science*, doi:10.1177/23998083251409526).

Consequence for the paper's AI claim. "AI turns city data into decisions" walks straight into Kitchin's fifth issue (instrumental rationality displacing deliberation) and into Huang's risk. The claim this paper makes instead is narrow and falsifiable: models lower the cost of *integration*, which is where the labour actually is; a model may narrate the record only behind a gate that admits nothing the record does not support; and the refusals are published as data. **The contribution is the gate and the record, not the model.**

3 · The inventory, re-read this evening

What	Verified value (9 Sept 2026, 13:52–15:05 UTC)
Source registry	211 records ; 14 publishers said no and are recorded with the clause
Permission evidence	188 sources with stored evidence · 313 captures · 169 access checks passed, 11 refused, 4 awaiting decision, 4 incomplete, 26 undocumented
Collectors polled	28 sources , cadence 10–60 min, each with a capture taken before the first read and re-checked weekly
Rows on disk / in the 24 h window	46 331 / 43 253
Hourly history	22 h , 843 series, from 2026-09-08T18 (public/history.json, new today)
Corrections	14 (C-001...C-014), appended, never edited
The mind	176 drops: 55 thoughts, 49 refused, 71 organelle, 1 retracted; 10 scored claims ; 124 thoughts in the published window, 11 voiced
Tests	123 offline tests green
Public site	rebuilt and pushed on a ~5 min cycle; four front-door answers in SR/EN/ZH/DE; a standing objection address

New since the morning: METAR at LYBE (S03) and the RHMZ river gauges (S52) as senses that publish their own measurement instant; the EDS day-after table (S54); the hourly history; the axonometric layer scheme with the legal frame as stratum zero; the related-work register; and the contact route.

4 • The contributions, restated against §2

1. **An audit of openness for one city, as a checkable register.** 211 access routes with a state each, 14 named refusals honoured, 26 undocumented admitted. This is Kitchin & McArdle's Dublin audit performed for Belgrade — the same finding (coverage is thinner than it looks) with the evidence attached rather than asserted.
2. **The metadata layer their §3 says is missing, actually built.** Every value carries three times, one of five states, a unit, and the captured permission that allowed it to be read; absence is a positive record with its reason; corrections are appended and never erased.
3. **Unknown measurement time as a first-class state, with a rendering grammar.** Parking publishes no measurement time; the row says so and the drawing smears its age rather than pretending to an instant. To our knowledge this is not standard practice on public city displays.
4. **A bounded mind, measured by its refusals.** Three small local models comment in Serbian; the validator admits only numbers present in the digest, requires citation inside the sentence, forbids a forecast stated as fact, and re-checks the language. 49 of 176 drops were refused and the refusals are public. This is an answer, in code, to Huang's risk — offered as a demonstration, not yet as an evaluation.
5. **A reflexive finding worth one paragraph of the paper (C-014).** The task that publishes this system's own record ran roughly a hundred times over seventeen hours reporting success — status Ready, last result 0, a log line per run — and published nothing, because a batch file that calls another batch file without `call` never returns. The failure was invisible precisely because the system's self-report was healthy. *An exit code of 0 is evidence that something finished, never evidence that it happened.* A paper about evidence should say that it caught itself doing exactly what it accuses dashboards of: trusting a receipt instead of the artefact.

5 • Recommended paper for 25 September

Title (working): *The instrument is not the bottleneck: an evidence-first observatory for Belgrade, and what it refuses to say.*

Abstract (≤ 200 words, draft). Cities are told they need more sensors. Belgrade already has a state air-quality network, meteorological and hydrological stations, airport observations, utility outage tables, parking counts and citizen sensors — and a resident still cannot establish what the air is doing nearby, how certain anyone is, and when it was measured. We argue the binding constraint is not sensing but the defensible record, and we test the argument by building one. A stdlib collector reads 28 permitted public sources on a 10–60 minute cadence; every value carries its source, three times, a unit, one of five evidential states and the captured permission that allowed it to be read; absence is recorded with its reason; corrections are appended and never erased. Of 211 registered access routes, 14 publishers refused and are honoured by name. Small local models comment on the record behind a validator that admits only what the record supports; 49 of 176 utterances were refused and the refusals are published. We report what the system cannot do, including a seventeen- hour failure it hid from itself, and argue that the contribution of AI here is the gate and the record rather than the model.

Six-page structure. (1) The gap, with §2's three anchors. (2) Method: permission before collection; the row and its three times; five states; absence. (3) The Belgrade audit: 211/28/14/26, with the table. (4) The bounded mind and its refusal counts. (5) C-014 and what self-reporting systems owe their readers. (6) Limits and what would falsify the claims.

6 • What must be true before submission

- ☐ The late-submission question actually asked (deadline 20 August passed; still unsent — this is the one blocking item and it is not a technical one).
 - ☐ Seven days of continuous history, so §3's table reports a week rather than 22 hours.
 - ☐ The gate evaluated, not only demonstrated: an adversarial set of digests and utterances with a false-accept rate. Without it, claim 4 stays a demonstration and the paper must say so.
 - ☐ The ISO 37120 mapping performed, or the claim in §4.1 explicitly scoped to our own categories.
 - ☐ Notification letters sent, so the paper can say the sources were told rather than merely not circumvented. Drafted; awaiting signature.
 - ☐ A rendering test in the repository, so the site's own defects stop being found by eye.
-

Honest verdict

- The thesis is defensible because its two load-bearing citations are hostile witnesses: the largest observatory in the comparison says insight is the bottleneck, and the standard critique of dashboards says the missing thing is precisely the layer we built. That is a much stronger position than the morning's inventory paper.
- The weakest sentence in the proposed abstract is "the contribution is the gate". It is currently supported by counted refusals over a few days, on one machine, with three small models — no baseline, no adversarial set, no inter-rater check on what a correct refusal is. A reviewer who presses there is right, and the paper must concede it in §6 rather than defend it.
- Twenty-two hours of history is thin for any claim about how values move over time. The honest move is to report the record's age everywhere it appears, which the site now does.
- Nothing in this work shows that better bookkeeping improves any actual decision. The most that is claimed is that a decision cannot be well made on a number whose age, origin and permission are unknown. A user study would be needed for more, and none has been run.
- The paper cannot be submitted at all until someone asks whether a late manuscript is accepted. That question has been open since the morning version and remains the only item on this list that no amount of work here can close.